

High Performance Computing: Tools and Applications

Edmond Chow
School of Computational Science and Engineering
Georgia Institute of Technology

Lecture 4

OpenMP directives

So far we have seen

```
#pragma omp parallel
```

```
#pragma omp for
```

Implied barrier at end of `for` construct

```
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i<n; i++)
    {
        a[i] = i;
    }

    // implied barrier

    // any thread sees all components of "a"
    // as updated
}
```

There is also an implied barrier at the end of sections and single constructs.

`firstprivate` and `lastprivate` clauses in `for` directive

- ▶ With the `private` clause, private variables are undefined at the beginning of the loop, and values within the loop are not visible after the loop
- ▶ `firstprivate` clause instead initializes the private variables
- ▶ `lastprivate` clause copies value of last iteration to the variable after the loop

Example: both firstprivate and lastprivate

```
#include <stdio.h>
#include <omp.h>

int main (int argc, char *argv[])
{
    int i, a = 1000;

    #pragma omp parallel for firstprivate(a) lastprivate(a)
    for (i=0; i<10; i++)
    {
        a = a + i;
    }
    printf("value of a: %d\n", a);

    return 0;
}
```

sections directive

- ▶ Used when different threads must execute different code
- ▶ Must still create threads with `parallel` directive
- ▶ In general, p threads created and n sections

```
#pragma omp parallel
#pragma omp sections
{
    #pragma omp section
    printf("First thread %d\n", omp_get_thread_num());

    #pragma omp section
    printf("Second thread thread %d\n", omp_get_thread_num());
}
```

What happens if $p < n$?

What happens if $p > n$?

Notes on sections directive

- ▶ at most n threads run in parallel
- ▶ can also use `firstprivate` and `lastprivate` with obvious definitions of *first* and *last*
- ▶ can combine `parallel` and `sections` directives `#pragma omp parallel sections` like `parallel for`

What is wrong with this code?

```
#pragma omp parallel
{
    a = 255;

    #pragma omp for
    for (i=0; i<n; i++)
        b[i] = a;
}
```


What is wrong with this code?

- ▶ Depending on hardware, write to `a` may not be atomic, and thread 0 may read `a` when thread 1 has only partially written to it.
- ▶ Possible solution is to use a barrier after the write.
- ▶ Multiple threads writing to `a` is also unnecessary.
- ▶ Solution here is to use `single` directive.

single directive

- ▶ Used when code should only be executed by a single thread
- ▶ Can be executed by any thread (related `master` directive)

```
#pragma omp parallel
{
    #pragma single
    a = 255;

    #pragma omp for
    for (i=0; i<n; i++)
        b[i] = a;
}
```

Implied barrier at end of structured block (of `single`).

- ▶ thread-safe update of shared variables
- ▶ generally requires the compiler to use atomic instructions in the instruction set
- ▶ applies to single statements only (not blocks) with specific forms of updating a memory location

```
#pragma omp atomic  
i = i + 1;
```

atomic directive: example allowed form

```
x = x binop expr;
```

where `x` is an l-value with scalar type, `expr` does not access the same storage as `x`, and `binop` is a binary operation, e.g., `+`.

For more details:

<https://software.intel.com/en-us/node/524509>

More OpenMP directives

```
#pragma omp sections
```

```
#pragma omp single
```

```
#pragma omp master
```

```
// no implied barrier on exit
```

```
#pragma omp barrier
```

```
#pragma omp ordered
```

```
// used inside parallel for loop
```

```
#pragma omp critical [name]
```

```
#pragma omp atomic
```

```
// only for statements of specific form
```

Some OpenMP clauses

- ▶ `num_threads` sets the number of threads in `parallel` directive
- ▶ `if` controls the `parallel` directive depending on a condition
- ▶ `nowait` removes the barrier at the end of `omp for` and other constructs
- ▶ `ordered` needed to indicate that an `ordered` directive is within an `omp for` loop

Example if clause

Only spawn threads if the “problem” is large enough:

```
#pragma omp parallel if (n > 1000)
```

More realistic particle simulations

- ▶ Particles have radius a
- ▶ Cubical simulation box has width L and *periodic* boundaries
- ▶ Particles interact with each other
 - ▶ repulsive force when they overlap
 - ▶ other forces, e.g., when particles are charged

Repulsive force when particles overlap

If distance s between particles i and j is less than $2a$, then force on particle i due to particle j is

$$f_{ij} = k_r(2a - s) \cdot \hat{n}$$

where $k_r = 100$ is the repulsion force constant and \hat{n} is the normal vector from j to i .

Updating the particle positions, including repulsive force

$$x(i) = x(i) + M \cdot f(i) \Delta t + \sqrt{2\Delta t} \cdot y(i)$$

where $f(i)$ is the total force on particle i and $M = 1$ is a constant.

Code to compute forces from the positions

```
for (i=0; i<np; i++) {  
    for (j=i+1; j<np; j++) {  
        ri = &pos[3*i];  
        rj = &pos[3*j];  
        dx = remainder(ri[0]-rj[0], L);  
        dy = remainder(ri[1]-rj[1], L);  
        dz = remainder(ri[2]-rj[2], L);  
  
        s2 = dx*dx + dy*dy + dz*dz;  
        if (s2 < 4.*a*a) {  
            s = sqrt(s2);  
            f = krepul*(2.-s);  
  
            forces[3*i+0] += f*dx/s;  
            forces[3*i+1] += f*dy/s;  
            forces[3*i+2] += f*dz/s;  
            forces[3*j+0] -= f*dx/s;  
            forces[3*j+1] -= f*dy/s;  
            forces[3*j+2] -= f*dz/s;  
        }  
    }  
}
```

How to parallelize this code using OpenMP?

Iterations on `i` are not independent, since different iterations can write to the same location in `forces`.

How to parallelize this code using OpenMP?

Iterations on `i` are not independent, since different iterations can write to the same location in `forces`.

- Update forces in a critical section

How to parallelize this code using OpenMP?

Iterations on `i` are not independent, since different iterations can write to the same location in `forces`.

- ▶ Update forces in a critical section
- ▶ Use atomic operations to update components of the force array

How to parallelize this code using OpenMP?

Iterations on i are not independent, since different iterations can write to the same location in `forces`.

- ▶ Update forces in a critical section
- ▶ Use atomic operations to update components of the force array
- ▶ Rewrite the outer loop so that iterations are independent (only update the forces for particle i , not j)

Exercise 4

- ▶ Update your code from Exercise 2 to include repulsive interactions between particles when they overlap. Parallelize using OpenMP and run on `jinx` or `deeptthought`.
- ▶ Compare the performance between using critical sections, atomic operations, and independent iterations (which do twice the number of distance computations).
- ▶ Submit your results in the `ex04` directory (do not forget to update your fork), including
 - ▶ `ex04.c` or `ex04.cpp` source file and `makefile`
 - ▶ `ex04.pdf` report with performance comparison
- ▶ *Due 10 pm, Monday, Sept. 5*